

Clustering Heart Disease Patient Data

Introduction

There are many industries where understanding how things group together is beneficial. For example, retailers want to understand the similarities among their customers to direct advertisement campaigns, and botanists classify plants based on their shared similar characteristics. One way to group objects is to use clustering algorithms. We will explore the usefulness of unsupervised clustering algorithms to help doctors understand which treatments might work with their patients. We are going to cluster anonymized data of patients who have been diagnosed with heart disease. Patients with similar characteristics might respond to the same treatments, and doctors could benefit from learning about the treatment outcomes of patients like those they are treating.

The dataset

The data we are analyzing comes from the V.A. Medical Center in Long Beach, CA. It has the following columns:

- id: patient identification number
- age: patient's age in years
- sex: patient gender(1 = male; 0 = female)
- cp: chest pain type(1 = typical angina, 2 = atypical angina, 3 = non-anginal pain, 4 = asymptomatic)
- trestbps: resting blood pressure(mm Hg)
- chol: serum cholesterol(mg/dl)
- fbs: fasting blood sugar > 120 mg/dl (1 = true, 0 = false)
- restecg: resting electrocardiograph results (0 = normal, 1 = having ST-T wave abnormality, 2 = showing probable or definite left ventricular hypertrophy)
- thalach: maximum heart rate achieved
- exang: exercise induced angina(1 = yes, 0 = no)
- oldpeak: ST depression induced by exercise relative to rest
- slope: the slope of the peak exercise ST segment

Let's load the datasets and packages and take a look at the first few rows.

```
```{r}
library(tidyverse)
library(ggplot2)

heart_disease <- read_csv("heart_disease_patients.csv")
```
```



```

```{r}
Scaling data and saving as a data frame
scaled <- scale(heart_disease)

Let's see the data after we scaled it
summary(scaled)
```

```

| age | sex | cp | trestbps |
|------------------|------------------|--------------------|-------------------|
| Min. : -2.8145 | Min. : -1.4549 | Min. : -2.2481 | Min. : -2.14149 |
| 1st Qu.: -0.7124 | 1st Qu.: -1.4549 | 1st Qu.: -0.1650 | 1st Qu.: -0.66420 |
| Median : 0.1727 | Median : 0.6851 | Median : -0.1650 | Median : -0.09601 |
| Mean : 0.0000 | Mean : 0.0000 | Mean : 0.0000 | Mean : 0.00000 |
| 3rd Qu.: 0.7259 | 3rd Qu.: 0.6851 | 3rd Qu.: 0.8765 | 3rd Qu.: 0.47218 |
| Max. : 2.4961 | Max. : 0.6851 | Max. : 0.8765 | Max. : 3.88132 |
| chol | fbs | restecg | thalach |
| Min. : -2.3310 | Min. : -0.4169 | Min. : -0.995103 | Min. : -3.4364 |
| 1st Qu.: -0.6894 | 1st Qu.: -0.4169 | 1st Qu.: -0.995103 | 1st Qu.: -0.7041 |
| Median : -0.1100 | Median : -0.4169 | Median : 0.009951 | Median : 0.1483 |
| Mean : 0.0000 | Mean : 0.0000 | Mean : 0.000000 | Mean : 0.0000 |
| 3rd Qu.: 0.5467 | 3rd Qu.: -0.4169 | 3rd Qu.: 1.015005 | 3rd Qu.: 0.7166 |
| Max. : 6.1283 | Max. : 2.3905 | Max. : 1.015005 | Max. : 2.2904 |
| exang | oldpeak | slope | |
| Min. : -0.6955 | Min. : -0.8954 | Min. : -0.9747 | |
| 1st Qu.: -0.6955 | 1st Qu.: -0.8954 | 1st Qu.: -0.9747 | |
| Median : -0.6955 | Median : -0.2064 | Median : 0.6480 | |
| Mean : 0.0000 | Mean : 0.0000 | Mean : 0.0000 | |
| 3rd Qu.: 1.4331 | 3rd Qu.: 0.4827 | 3rd Qu.: 0.6480 | |
| Max. : 1.4331 | Max. : 4.4445 | Max. : 2.2708 | |

Empirical Analysis

Now that we have scaled the data, we can start the clustering process. For the k-means algorithm, it is necessary to select the number of clusters in advance. It is also essential to make sure that our results are reproducible when conducting a statistical analysis. This means that they will get the same results when someone runs our code on the same data. Because we are doing an analysis with a random aspect, it is necessary to set seed to ensure reproducibility. Reproducibility is essential because doctors will potentially use our results to treat patients. Other analysts must see where the groups come from and can verify the results.

```

```{r}
Set the seed so that results are reproducible
set.seed(10)

Select a number of clusters
k <- 5

```

```
Run the k-means algorithm
first_clust <- kmeans(scaled, centers = k, nstart = 1)
```

```
How many patients are in each cluster?
first_clust$size
```
```

```
66 43 88 61 45
```

Because the k-means algorithm initially selects the cluster centers by randomly selecting points, different algorithm iterations can result in different clusters. If the algorithm genuinely groups similar observations (as opposed to clustering noise), then cluster assignments will be somewhat robust between various iterations of the algorithm.

Regarding the heart disease data, this would mean that the same patients would be grouped even when the algorithm is initialized at different random points. If patients are not in similar clusters with various algorithm runs, the clustering method does not pick up on meaningful relationships between patients.

We're going to explore how the patients are grouped with another iteration of the k-means algorithm. We will then be able to compare the resulting groups of patients.

```
``{r}
# Set the seed
set.seed(38)

# Select a number of clusters and run the k-means algorithm
k <- 5
second_clust <- kmeans(scaled, centers = k, nstart = 1)
```

```
# How many patients are in each cluster?
second_clust$size
```
```

```
65 43 61 46 88
```

The clusters must be stable. Even though the algorithm begins by randomly initializing the cluster centers, if the k-means algorithm is the right choice for the data, then different initializations of the algorithm will result in similar clusters.

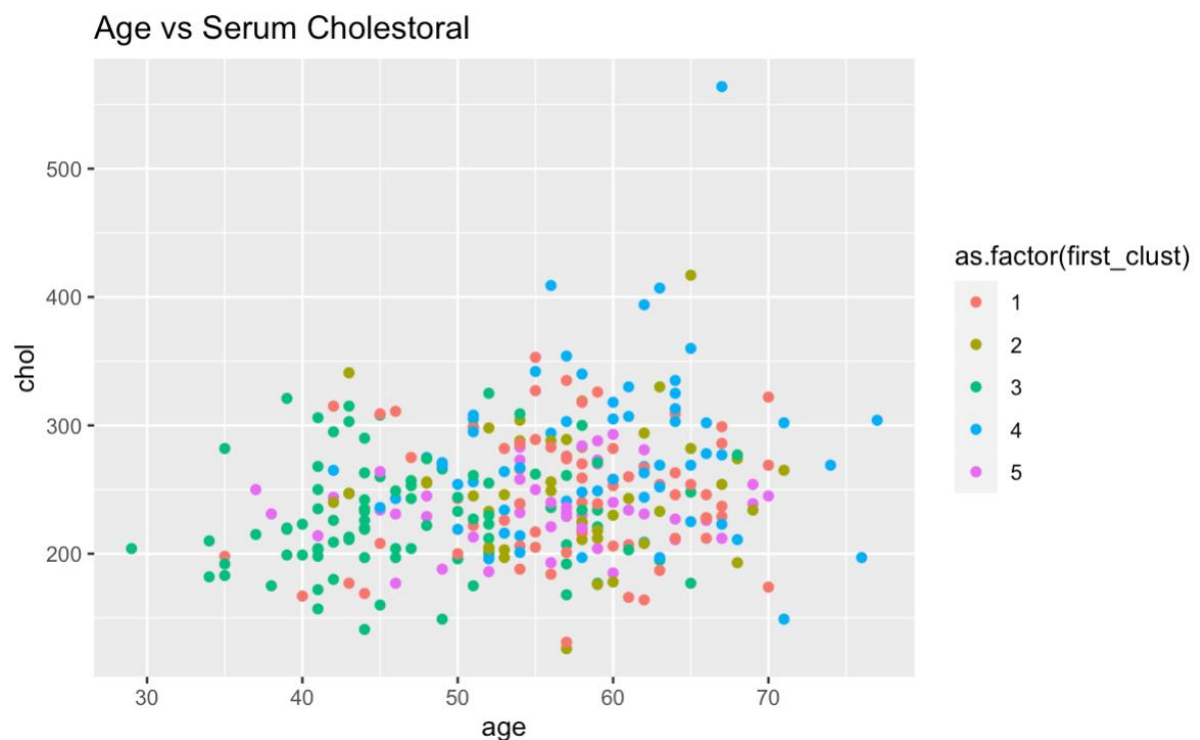
The clusters from different iterations may not be the same, but the clusters should be roughly the same size and have similar distributions of variables. If there is a lot of change in clusters between different algorithm iterations, then k-means clustering is not the right choice for the data.

It is impossible to validate that the clusters obtained from the algorithm are accurate because there is no patient labeling. Thus, it is necessary to examine how the clusters


change between different iterations of the algorithm. We're going to use some visualizations to get an idea of the cluster stabilities. That way, we can see how specific patient characteristics may have been used to group patients together.

```
``{r}
Add cluster assignments to the data
heart_disease["first_clust"] <- first_clust$cluster
heart_disease["second_clust"] <- second_clust$cluster

Create and print the plot of age and chol for the first clustering algorithm
plot_one <- ggplot(heart_disease, aes(x=age, y=chol, color=as.factor(first_clust))) +
 geom_point() + labs(title = "Age vs Serum Cholestorl")
plot_one
``
```



```
``{r}
Create and print the plot of age and chol for the second clustering algorithm
plot_two <- ggplot(heart_disease, aes(x=age, y=chol, color=as.factor(second_clust))) +
 geom_point() + labs(title = "Age vs Serum Cholestorl")
plot_two
``
```



An alternative to k-means clustering is hierarchical clustering. This method works well when data have a nested structure. Heart disease patient data might follow this type of structure. For example, if men are more likely to exhibit specific characteristics, those characteristics might be nested inside the gender variable. Hierarchical clustering also does not require the number of clusters to be selected before running the algorithm.


Clusters can be selected by using the dendrogram. The dendrogram allows us to see how similar observations are to one another, and they are useful in helping us choose the number of clusters to group the data. It is now time for us to see how hierarchical clustering groups the data.

```
```{r}
# Execute hierarchical clustering with complete linkage
hier_clust_1 <- hclust(dist(scaled), method = "complete")

# Print the dendrogram
plot(hier_clust_1)

# Get cluster assignments based on number of selected cluster
hc_1_assign <- cutree(hier_clust_1, 5)
```
```





The doctors are interested in grouping similar patients together to determine appropriate treatments. Therefore, they want clusters with more than a few patients to see different treatment options. While a patient can be in a cluster by themselves, the treatment they received might not be recommended for someone else in the group.

Like the k-means algorithm, the way to evaluate hierarchical clusters is to investigate which patients are grouped together. We're going to examine the clusters resulting from the two hierarchical algorithms.

```
``{r}
```

```
Add assignment of chosen hierarchical linkage
```

```
heart_disease["hc_clust"] <- hc_1_assign
```

```
Remove the sex, first_clust, and second_clust variables
```

```
hd_simple <- heart_disease[, ! (names(heart_disease) %in% c("sex", "first_clust",
"second_clust"))]
```

```
Get the mean and standard deviation summary statistics
```

```
clust_summary <- do.call(data.frame, aggregate(. ~hc_clust, data = hd_simple, function(x)c
(avg = mean(x), sd = sd(x))))
```

```
clust_summary
```

```
``
```



| hc_clust<br><int> | age.avg<br><dbl> | age.sd<br><dbl> | cp.avg<br><dbl> | cp.sd<br><dbl> | trestbps.avg<br><dbl> | trestbps.sd<br><dbl> |
|-------------------|------------------|-----------------|-----------------|----------------|-----------------------|----------------------|
| 1                 | 51.41667         | 8.540979        | 2.783333        | 0.9470625      | 129.1389              | 15.93800             |
| 2                 | 58.11111         | 7.754246        | 3.763889        | 0.6165112      | 130.0417              | 13.90657             |
| 3                 | 61.00000         | 3.908034        | 3.916667        | 0.2886751      | 168.5000              | 17.45904             |
| 4                 | 59.00000         | 9.203580        | 3.571429        | 0.8501112      | 134.7714              | 18.64070             |
| 5                 | 64.75000         | 2.061553        | 3.250000        | 0.5000000      | 138.7500              | 18.42779             |

5 rows | 1-7 of 21 columns

| chol.avg<br><dbl> | chol.sd<br><dbl> | fbs.avg<br><dbl> | fbs.sd<br><dbl> | restecg.avg<br><dbl> | restecg.sd<br><dbl> |
|-------------------|------------------|------------------|-----------------|----------------------|---------------------|
| 239.8722          | 42.29228         | 0.1222222        | 0.3284559       | 0.8444444            | 0.9905826           |
| 253.2222          | 49.74476         | 0.1805556        | 0.3873488       | 1.4027778            | 0.9140488           |
| 284.9167          | 53.00336         | 0.3333333        | 0.4923660       | 1.2500000            | 0.9653073           |
| 233.8571          | 49.67136         | 0.1428571        | 0.3550358       | 0.6857143            | 0.9321521           |
| 433.7500          | 89.93470         | 0.2500000        | 0.5000000       | 2.0000000            | 0.0000000           |

5 rows | 8-13 of 21 columns

| thalach.avg<br><dbl> | thalach.sd<br><dbl> | exang.avg<br><dbl> | exang.sd<br><dbl> | oldpeak.avg<br><dbl> | oldpeak.sd<br><dbl> |
|----------------------|---------------------|--------------------|-------------------|----------------------|---------------------|
| 161.5722             | 15.779214           | 0.07777778         | 0.2685686         | 0.555000             | 0.7847196           |
| 135.5417             | 17.991342           | 0.81944444         | 0.3873488         | 1.451389             | 1.0804268           |
| 147.7500             | 13.157266           | 0.75000000         | 0.4522670         | 2.316667             | 1.4708274           |
| 116.8857             | 17.842071           | 0.48571429         | 0.5070926         | 2.240000             | 1.3856831           |
| 156.2500             | 3.774917            | 0.00000000         | 0.0000000         | 1.100000             | 0.3829708           |

5 rows | 14-19 of 21 columns

| thalach.sd<br><dbl> | exang.avg<br><dbl> | exang.sd<br><dbl> | oldpeak.avg<br><dbl> | oldpeak.sd<br><dbl> | slope.avg<br><dbl> | slope.sd<br><dbl> |
|---------------------|--------------------|-------------------|----------------------|---------------------|--------------------|-------------------|
| 15.779214           | 0.07777778         | 0.2685686         | 0.555000             | 0.7847196           | 1.388889           | 0.5730336         |
| 17.991342           | 0.81944444         | 0.3873488         | 1.451389             | 1.0804268           | 1.750000           | 0.5240686         |
| 13.157266           | 0.75000000         | 0.4522670         | 2.316667             | 1.4708274           | 2.166667           | 0.5773503         |
| 17.842071           | 0.48571429         | 0.5070926         | 2.240000             | 1.3856831           | 2.200000           | 0.4058397         |
| 3.774917            | 0.00000000         | 0.0000000         | 1.100000             | 0.3829708           | 1.500000           | 0.5773503         |

5 rows | 15-21 of 21 columns


In addition to looking at the distributions of variables in each hierarchical clustering run, we will make visualizations to evaluate the algorithms. Even though the data has more than two dimensions, we can understand how the data clusters by looking at a scatter plot of two variables. We want to look for patterns that appear in the data and see what patients get clustered together.

```

```{r}
# Plot age and chol
plot_one <- ggplot(heart_disease, aes(x=age, y = chol, color =as.factor(hc_clust))) +
  geom_point() + labs(title = "Age vs Serum Cholesterol")
plot_one
```

```

Age vs Serum Cholesterol




```

{r}
Plot oldpeak and trestbps
plot_two <- ggplot(heart_disease, aes(x = oldpeak, y = trestbps, color = as.factor(hc_clust)))
+ geom_point() + labs(title = "ST Depression vs Resting Blood Pressure")
plot_two

```

ST Depression vs Resting Blood Pressure




## Conclusion

Now that we've tried multiple clustering algorithms, it is necessary to determine if we think any of them will work for clustering our patients. For the k-means algorithm, similar clusters must be produced for each algorithm iteration to make sure that the algorithm clusters the signal, not the noise.

It's also important for the k-mean algorithm to seem stable when running multiple iterations. This means that we would see similar groups of patients showing up in the plots from the different iterations of the algorithm. For the hierarchical clustering, we need a method that puts a balanced number of patients in each group.

So the best algorithm that shows promise for this data is Hierarchical clustering with complete linkage.



## Source:

- <https://archive.ics.uci.edu/ml/datasets/heart+Disease>